

7. Scaled Dot-Product Attention (Transformer)

Core mechanism of the MarianMT encoder-decoder architecture.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^{\top}}{\sqrt{d_k}}\right)V$$

Q = queries, K = keys, V = values, d_k = key dimension