

9. Beam Search Decoding

Selects the translation with highest log-probability across beams.

$$\hat{Y} = \arg \max_{Y \in \mathcal{B}} \sum_{t=1}^{|Y|} \log P_{\theta}(y_t \mid y_{<t}, X)$$

$\mathcal{B}$ = set of candidate sequences (beam width = 4 in our experiment)